

# The Dual-MINDy/gBPKF Package

Matthew Singh

January 22, 2025

## Introduction

This MATLAB package provides the tools to implement the gBPKF algorithm as introduced in the papers:

Efficient identification for modeling high-dimensional brain dynamics (2022).

Precision data-driven modeling of cortical dynamics reveals person-specific mechanisms underpinning brain electrophysiology (2025).

The documentation is a work-in-progress and more details will be added. This also doubles as an extended toolbox for training vanilla-RNNs with standard-BPTT (no Kalman Filter) being a very limited subset of its functionality. Thus, there are many features and options well-outside the range used in the above papers (e.g. many gradient-optimizers besides NADAM).

This code seeks to solve problems of the form:

$$x_{t+1} = f_{\theta}(x_t) + \varepsilon_t \tag{1}$$

$$y_t = H_t x_t + \eta_t \tag{2}$$

In terms of parameters  $\theta$ . Here,  $x_t$  is the unknown, true state of the system, evolving according to  $f$  and process noise  $\varepsilon$  with covariance  $Q_t$ . The observations  $y_t$  are a linear transformation of  $x_t$  defined by the measurement matrix  $H_t$  with the addition of measurement noise  $\eta_t$  with covariance  $R_t$ . You will note that

in our formulation, all components (except)  $\theta$  are potentially time-varying. In the current implementation, the measurement transformation  $H_t$  and the noise covariances  $Q_t$ ,  $R_t$  are allowed to differ between user-defined segments of data (e.g. different recordings) but are not currently coded to vary continuously.

The objective is to recover the model parameters  $\theta$  (and potentially  $Q_t$ ) using only the observations  $y_t$ . The heart of the gBPKF algorithm involves further estimating the true system-states  $x_t$  for small chunks of data. However, the current implementation does not return a final estimate of  $x_t$ .

In this implementation, we start with the generic RNN form for the process:

$$f(x_t) = W\psi(x) + Dx_t + c \quad (3)$$

$$\psi(x) := \tanh(s \circ x + v) \quad (4)$$

Hence, the unknown parameters consist of the connection matrix  $W$ , the gain  $s$ , the shift  $v$ , the discrete time-constant  $D$ , the baseline "drive" vector  $c$ .

This form allows a vast variety of biological models by placing additional constraints on the parameters.

For instance, the EI models used in the original two papers can be formed by structuring  $x_t$  to have both excitatory and inhibitory activation. For 100 regions we might choose the first 100 elements of  $x_t$  to be excitatory activity and the next 100 to be inhibitory activity, so that  $x_t$  is 200-dimensional. We add meaning to neuron-types through constraints on  $W$ . In this case,  $W$  has the block form (with each block  $100 \times 100$ ):

$$W := \begin{bmatrix} W^{EE} & \beta^{IE} \\ W^{EI} & \beta^{II} \end{bmatrix} \quad (5)$$

Here, excitatory connection matrices ( $W^{EE, EI}$ ) are constrained to be non-negative. By contrast, inhibitory connections  $\beta^{II, IE}$  are constrained to be non-positive and diagonal, reflecting the local nature of interneurons. The corresponding measurement matrix is

$$H = \begin{bmatrix} -L & 0_{m \times n} \end{bmatrix} \quad (6)$$

With  $L$  denoting the leadfield-matrix for channel-level data or (negative) identity for source-localized data (i.e.  $-L = I_n$ ). The zeros matrix on the right, reflects that inhibitory cells do not generally contribute to the MEG/EEG signal.

Throughout we make use of the term "population" to denote one element of the state-vector  $x_t$ . In the original papers there are 100 brain regions with 200 total populations: one excitatory and one inhibitory per region.

Note that  $EI$  is read as the connection from excitatory to inhibitory. Within matrices  $W_{a,b}$  is read as the connection to  $a$  from  $b$ . Model-specification generally should contain additional constraints as we discuss later.

We also strongly recommend constraining  $W^{EE,EI}$  with a mask specifying which connections are forced-zero as demonstrated in the next section. Unfortunately, brain folding generates colinearities (especially around the midline) that can't be resolved without further constraints.

A full simulation example with documentation is provided in the script: `BPKF_Example_Sim`

## 1 General Remarks

Users will note that the package contains a very large number of parameter options. Many of these were never changed and are unlikely to affect results (i.e. should just be close to zero or one), but we offer full control.

The gBPKF algorithm is implemented in the main function `BPKF_Full`. This function takes 6 arguments (a 7th argument for inputs is under-construction).

The initial `Xguess` argument (for an initial guess of  $x_t$  is currently deprecated, as the recommended methods do not need an initial guess. For continuity's sake, however, we are still requiring this to be passed.

`Xguess`: The initial guess of  $x_t$  (can be random, this is **deprecated** for now)

`MeasSet`: The measurements  $y_t$

`ParStr`: General settings

`KalSpec`: Kalman-Filter settings

`ModelSpec`: Model-Specification

`GradSpec`: Gradient-optimizer settings

## 2 Data Formatting

### 2.1 Xguess (deprecated) and MeasSet

The first two arguments are cell-arrays of timeseries. Each element of the cell-array is one continuous recording specified as an  $n \times t$  matrix for  $n$  variables and  $t$  timepoints per recording. The arguments Xguess and MeasSet should contain the same number of matrices (recordings). Recordings can be different lengths and contain different channel configurations/counts (see later, not recommended for beginners). However, for each recording, the number of timepoints  $t$  should be the same between the measurements and the initial state-guess. Because we are fitting a single model, the number of rows in each Xguess matrix should be the same and equal to the dimension of  $x_t$ . Small random values are fine for Xguess which is currently deprecated (ignore it). The number of columns for each matrix in MeasSet should match the number of channels for that recording.

## 3 ParStr

The ParStr structure specifies some high-level model options including the loss function, batch size, number of batches, and settings for recording error, intermediary estimates, etc.

### 3.1 Configuring Data Sampling

At each iteration (batch), one-or-more data-chunks (minibatches) are selected as training-data. Gradients are accumulated during each minibatch and the model is updated at each batch iteration. These functions are controlled with:

ParStr.BatchSz = number of minibatches per batch  
ParStr.NBatch = number of batch iterations

When using EKF estimation (recommended), several proximal time-segments are selected per minibatched and share a common covariance estimate. The number of time-segments is specified by ParStr.nStack. By default, these time-segments are spaced 5 time-points apart, but this value can be changed by setting ParStr.ShiftSpace equal the desired value. The use of overlapping time-

segments allows efficient memory access while retaining some sensitivity to non-stationary features of the covariance (see SI in the 2025 paper).

## 4 KalSpec

KalSpec specifies the Kalman-Filter implementation. In brief, the forward-pass of the gBPKF algorithm contains 3 phases: 1) the initial estimate of  $x_0, P_0$ ; 2) refining  $x_t$  estimates using the Kalman Filter; and 3) free-running multi-step predictions (without filtering) once a good  $x_t$  is found.

### 4.1 Specifying the Initial Estimate

The initial distribution estimates, based on mean-square-error are fully determined by  $y_0$  and the covariances of  $x$  and  $y$ . Replacing  $cov(y)$  with its theoretical value  $(H_t cov(x) H_t^T + R_t)$  further yields an equation in terms of  $y_0$  (known) and  $cov(x)$  (unknown). The code includes two ways to estimate  $cov(x)$ : through simulating solutions (recommended) or by directly learning  $cov(x)$  (not recommended).

The recommended way (simulation) is engaged by setting `KalSpec.BFcov='s'`. The simulation size is dictated by the fields: `KalSpec.nSim`: number of simulations to run (in parallel) `KalSpec.SimLength`: simulation duration

In general, long simulations are required to estimate a high-dimensional covariance matrix. We speed things up in two ways: 1) For efficiency, we reuse the end of previous simulations to randomly re-seed initial conditions of new simulations. This enables fairly short simulations per-batch as we don't have to wait for the system to approach the steady-state distribution. The number of previous values saved for reseeding is set by `KalSpec.nSaveStart`. 2) The covariance estimate  $\Sigma_0$  is autoregressively updated between batches with AR-coefficient set by `KalSpec.decP=u`  $u \in (0, 1)$ . For batch  $j$ :

$$\Sigma_0^{j+1} = (1 - u)cov(x_{sim}) + u * \Sigma_0^j \quad (7)$$

This estimate is useful, by the end of model-training but can be very ill-behaved when the model is random (early training). Therefore, we use a convex combination of this estimate and a prior, baseline estimate of  $cov(x)$ :

$$cov(x) \approx (1 - q)\Sigma_0 + q * P_{fix} \quad (8)$$

Here, the weighting constant  $q$  is set by `KalSpec.decFix= $q \in (0,1)$`  and the baseline covariance is set by `KalSpec.Pfix= $P_{fix}$` . For MEG data, we used:

```
KalSpec.BFcov='s';
KalSpec.Pfix=eye(n)/4;
KalSpec.decFix=0.1;
KalSpec.decP=0.95;
KalSpec.SimLength=50;
KalSpec.nSim=50;
KalSpec.nSaveStart=250;
```

Changing these settings (if desired) ultimately comes down to balancing run-time, numerical-stability, and accuracy-per-iteration.

Alternatively (not recommended) one can treat  $cov(x)$  as an unknown parameter with the functional form:  $cov(x) = P_{fix} + P_{rt}P_{rt}^T$ . Here  $P_{fix}$  is fixed, positive-definite and  $P_{rt}$  is learned along with the other parameters. This option is specified by setting `KalSpec.BF='y'`; `KalSpec.Pfix= $P_{fix}$` . Finally, the initial value of  $P_{rt}$  is specified (via Cholesky decomposition) by setting `KalSpec.Pbase= $P_{rt}P_{rt}^T$` .

## 4.2 Specifying the loss-function

All loss functions take the form:

$$H(z^T M z) \tag{9}$$

in which  $H$  is a positive scalar function,  $M$  is a positive semi-definite matrix, and  $z$  is the model-prediction error in terms of predicting  $y$ . By default, the loss function is quadratic so  $H(x)=x$ , but  $H$  can be set to Huber loss by specifying `ParStr.CostFun='Huber'`;

If using Huber-loss, the fields `ParStr.decHuber` and `ParStr.HuberScale` should be specified. In this case, the Huber-loss parameter  $\alpha$  updates according to

$$\alpha_{t+1} = d\alpha_t + (1 - d)c * med(L) \tag{10}$$

In which `ParStr.decHuber= $d$` , `ParStr.HuberScale= $c$`  and  $L$  denotes the median quadratic loss  $(z^T M z)$ -across samples for the current minibatch. Values of  $\alpha$  are specific to recording, and prediction-step. This is meant to auto-adjust  $\alpha$  across setups. In the 2025 paper, we used: `ParStr.decHuber=0.99`; `ParStr.HuberScale=2`;

The cost-matrix  $M$  is specified through `ParStr.ErrMat`.  $M$  can be specified as a cell-array of matrices (one-per recording) or a single matrix if all recordings have the same number of channels. You can pass specific values (a matrix or cell of matrices). For convenience, this field also accepts the string 'L2' which makes  $M := I_n$ , or the string 'Mahal' which makes  $M_t := (H_t Q_t H_t^T + R_t)^{-1}$ . This quantity is the Mahalanobis Distance when  $x_{t-1}$  is known with certainty.